

# Lecture 8 - Reasoning under Uncertainty (Part II)

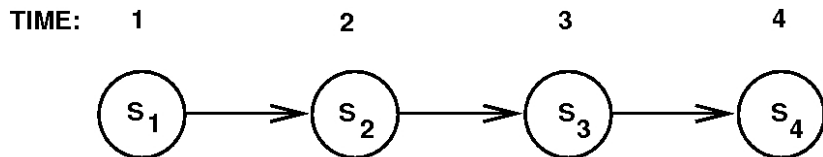
Jesse Hoey  
School of Computer Science  
University of Waterloo

June 13, 2022

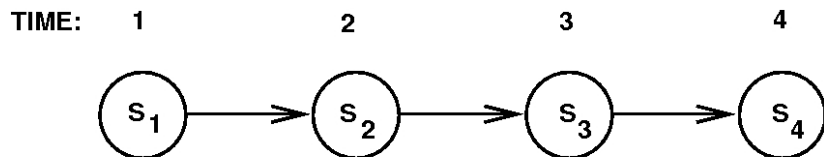
Readings: Poole & Mackworth (2nd ed.) Chapt. 8.5 - 8.9

# Probability and Time

- A node repeats over **time**
- **explicit** encoding of time
- chain has length = amount of time you want to model
- **event-driven** times or **clock-driven** times
- e.g. **Markov chain**



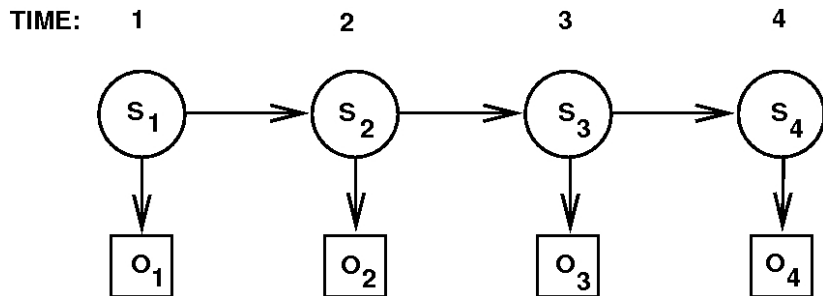
# Markov assumption



$$P(S_{t+1}|S_1, \dots, S_t) = P(S_{t+1}|S_t)$$

This distribution gives the **dynamics** of the Markov chain

# Hidden Markov Models (HMMs)



Add: observations  $O_t$  (always observed, so the node is square) and

observation function  $P(O_t|S_t)$

Given a sequence of observations  $O_1, \dots, O_t$ , can estimate

**filtering**:

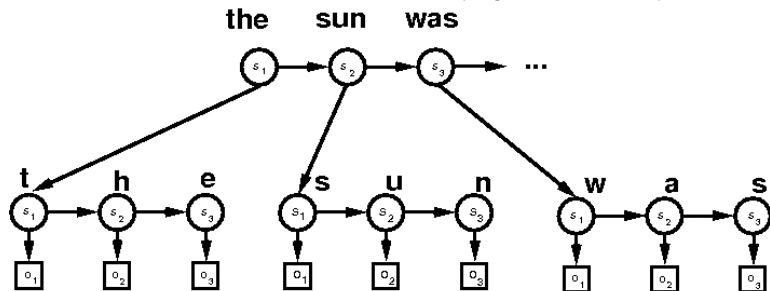
$$P(S_t|O_1, \dots, O_t)$$

or **smoothing**, for  $k < t$

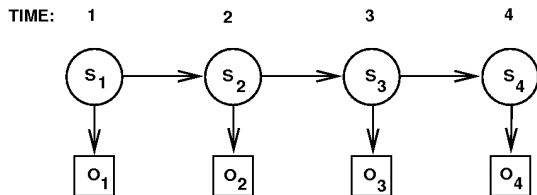
$$P(S_k|O_1, \dots, O_t)$$

# Speech Recognition

- Most well known application of HMMs
- **observations** : audio features
- **states** : phonemes
- **dynamics** : models e.g. co-articulation
- **HMMs** : words
- Can build hierarchical models (e.g. sentences)



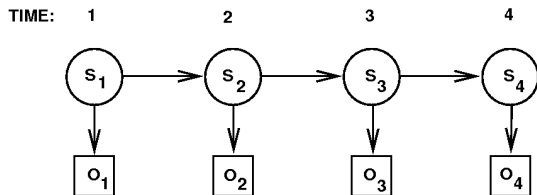
# Belief Monitoring in HMMs



filtering:

$$\begin{aligned}\alpha_i &= P(S_i | o_0 \dots, o_i) \\ &\propto P(S_i, o_0, \dots, o_i) \\ &= P(o_i | S_i) \sum_{S_{i-1}} P(S_i, S_{i-1}, o_0, \dots, o_{i-1}) \\ &= P(o_i | S_i) \sum_{S_{i-1}} P(S_i | S_{i-1}) P(S_{i-1}, o_0, \dots, o_{i-1}) \\ &\propto P(o_i | S_i) \sum_{S_{i-1}} P(S_i | S_{i-1}) \alpha_{i-1}\end{aligned}$$

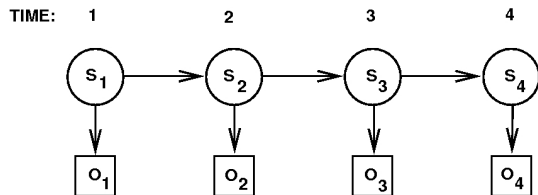
# Belief Monitoring in HMMs



smoothing:

$$\begin{aligned}\beta_{i+1} &= P(o_{i+1} \dots, o_T | S_i) \\ &= \sum_{S_{i+1}} P(S_{i+1}, o_{i+1}, \dots, o_T | S_i) \\ &= \sum_{S_{i+1}} P(o_{i+1} | S_{i+1}, o_{i+2}, \dots, o_T, S_i) P(S_{i+1}, o_{i+2}, \dots, o_T | S_i) \\ &= \sum_{S_{i+1}} P(o_{i+1} | S_{i+1}) P(o_{i+2}, \dots, o_T | S_{i+1}, S_i) P(S_{i+1} | S_i) \\ &= \sum_{S_{i+1}} P(o_{i+1} | S_{i+1}) P(S_{i+1} | S_i) \beta_{i+2}\end{aligned}$$

# Belief Monitoring in HMMs



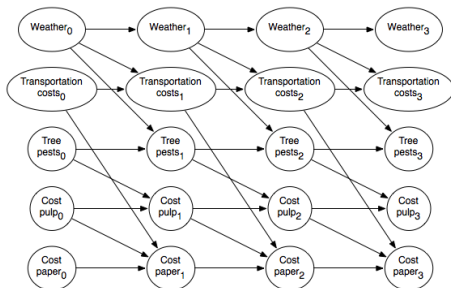
filtering and smoothing together :

$$\alpha_i \beta_{i+1} = P(o_{i+1} \dots, o_T | S_i) P(S_i | o_0 \dots, o_i) \propto P(S_i | O)$$



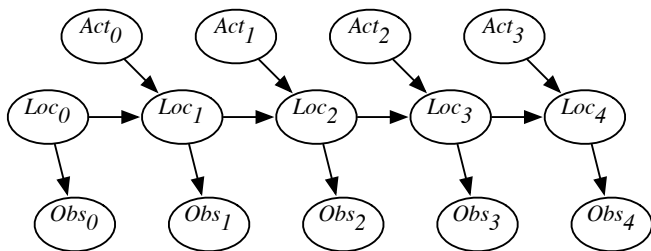
# Dynamic Bayesian Networks (DBNs)

- in general, any Bayesian network can repeat over time:  
**DBN**
- Many examples can be solved with **variable elimination**,
- may become too complex with enough variables
- **event-driven** times or **clock-driven** times



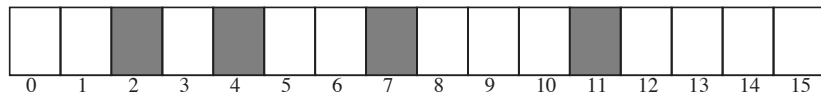
# Example: localization

- Suppose a robot wants to determine its location based on its actions and its sensor readings: **Localization**
- This can be represented by the **augmented HMM**:



# Example localization domain

- **Circular** corridor, with 16 locations:



- Doors at positions: 2, 4, 7, 11.
- Noisy Sensors
- Stochastic Dynamics
- Robot starts at an unknown location and must determine where it is, known as the **kidnapped robot** problem.
- see handout `robotloc.pdf`

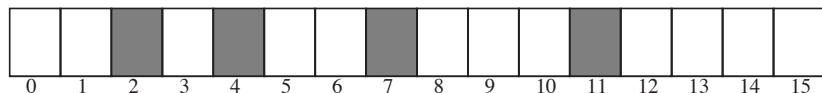
# Example Sensor Model

- $P(\text{Observe Door} \mid \text{At Door}) = 0.8$
- $P(\text{Observe Door} \mid \text{Not At Door}) = 0.1$

# Example Dynamics Model

- $P(Loc_{t+1} = l | Action_t = goRight \wedge Loc_t = l) = 0.1$
- $P(Loc_{t+1} = l + 1 | Action_t = goRight \wedge Loc_t = l) = 0.8$
- $P(Loc_{t+1} = l + 2 | Action_t = goRight \wedge Loc_t = l) = 0.074$
- $P(Loc_{t+1} = l' | Action_t = goRight \wedge Loc_t = l) = 0.002$  for any other location  $l'$ .
  - ▶ All location arithmetic is modulo 16.
  - ▶ The action *goLeft* works the same but to the left.

## Example sequence

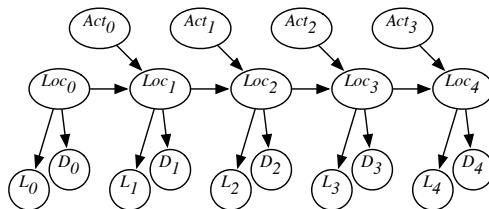


observe door, go right, observe no door, go right, observe door  
where is the robot?

$$P(\text{Loc}_2 = 4 | O_0 = d, A_0 = r, O_1 = \neg d, A_1 = r, O_2 = d) = 0.42$$

# Combining sensor information

- **Example:** we can combine information from a light sensor and the door sensor **Sensor Fusion**
- **Key Point:** Bayesian probability ensures that evidence is integrated **proportionally to its precision.**
- Sensors are **precision weighted**



$Loc_t$  robot location at time  $t$

$D_t$  door sensor value at time  $t$

$L_t$  light sensor value at time  $t$

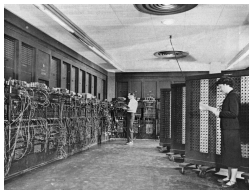
# Probability Distribution and Monte Carlo



John von Neumann  
1903 - 1957



Stanislaw Ulam  
1909-1984



ENIAC  
1949



Monte Carlo  
1949



# Stochastic Simulation

- **Idea:** probabilities  $\leftrightarrow$  samples
- Get probabilities from samples:

$X$	<i>count</i>
$x_1$	$n_1$
$\vdots$	$\vdots$
$x_k$	$n_k$
<i>total</i>	$m$

 $\leftrightarrow$ 

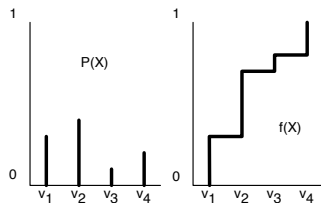
$X$	<i>probability</i>
$x_1$	$n_1/m$
$\vdots$	$\vdots$
$x_k$	$n_k/m$

- If we could **sample** from a variable's (posterior) probability, we could **estimate** its (posterior) probability.

# Generating samples from a distribution

For a variable  $X$  with a discrete domain or a (one-dimensional) real domain:

- **Totally order** the values of the domain of  $X$ .
- Generate the **cumulative probability distribution** :  
 $f(x) = P(X \leq x)$ .
- Select a value  $y$  **uniformly** in the range  $[0, 1]$ .
- Select the  $x$  such that  $f(x) = y$ .



# Hoeffding Bound

$p$  is true probability,  $s$  is sample average,  $n$  is number of samples

- $P(|s - p| > \epsilon) \leq 2e^{-2n\epsilon^2}$
- if we want an error greater than  $\epsilon$  in less than a fraction  $\delta$  of the cases, solve for  $n$ :

$$2e^{-2n\epsilon^2} < \delta$$

$$n > \frac{-\ln \frac{\delta}{2}}{2\epsilon^2}$$

- we have

$\epsilon$ error	cases with error $> \epsilon$	samples needed
0.1	1/20	184
0.01	1/20	18,445
0.1	1/100	265

# Forward sampling in a belief network

- Sample the variables **one at a time** ;
- sample **parents** of  $X$  **before** you sample  $X$ .
- Given values for the parents of  $X$ , sample from the **probability of  $X$  given its parents** .
- for samples  $s_i, i = 1 \dots N$ :

$$P(X = x_i) \propto \sum_{s_i} \delta(x_i) = N_{X=x_i}$$

where

$$\delta(x_i) = \begin{cases} 1 & \text{if } X = x_i \text{ in } s_i \\ 0 & \text{otherwise} \end{cases}$$

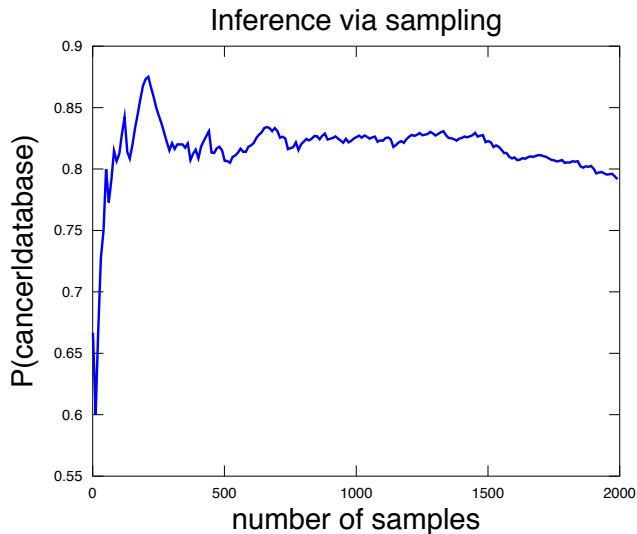
# Sampling for a belief network: inference

Sample	Malfnction	Cancer	TestB	TestA	Report	Database
$s_1$	false	false	true	true	false	false
$s_2$	false	true	true	true	true	true
$s_3$	false	true	true	true	true	true
$s_4$	false	false	false	true	false	false
$s_5$	true	true	true	true	false	false
$s_6$	false	true	false	true	false	false
$s_7$	false	false	false	true	false	true
...						
$s_{1000}$	false	false	false	true	false	false

To get  $P(H = h_i | E = e_i)$  simply

- count the number of samples that have  $H = h_i$  and  $E = e_i$ ,  $N(h_i, e_i)$
- divide by the number of samples that have  $E = e_i$ ,  $N(e_i)$
- $P(H = h_i | E = e_i) = \frac{P(H=h_i \wedge E=e_i)}{P(E=e_i)} = \frac{N(h_i, e_i)}{N(e_i)}$
- $P(C = \text{True} | \text{Database} = \text{True})$  based on first 7 samples?

# Forward Sampling

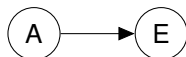


# Rejection Sampling

- To estimate a posterior probability given evidence  $Y_1 = v_1 \wedge \dots \wedge Y_j = v_j$ :
- If, for any  $i$ , a sample assigns  $Y_i$  to any value other than  $v_i$  **reject that sample**.
- The **non-rejected** samples are distributed according to the posterior probability.
- in the Hoeffding bound,  $n$  is the number of **non-rejected samples**

# Example Network

$P(A = \text{true})$
0.4



$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

If we draw  $N$  samples  $s_{i=1\dots N}$  by

- sampling  $A$ :  $a_{i=1\dots N}$
- sampling from  $E$  given  $A$ :  $e_{i=1\dots N}$

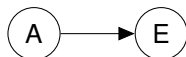
then

- $\approx N_t = 0.4N$  of them will have  $A = \text{true}$ , and of these  $\approx 10\%$  will have  $E = \text{true}$
- $\approx N_f = 0.6N$  of them will have  $A = \text{false}$ , and of these  $\approx 30\%$  will have  $E = \text{true}$



# Example Network

$$\frac{P(A = \text{true})}{0.4}$$



$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

so we have

$A$	$E$	$N_{AE}$
true	false	$N_{tf} = 0.4 \times 0.9 \times N$
true	true	$N_{tt} = 0.4 \times 0.1 \times N$
false	false	$N_{ff} = 0.6 \times 0.7 \times N$
false	true	$N_{ft} = 0.6 \times 0.3 \times N$

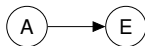
We want to compute

$$P(a|e) = P(A = \text{true} | E = \text{true}) \propto \sum_{s_i} \delta(a_i = \text{true}) \delta(e_i = \text{true})$$

$$\begin{aligned} P(a|e) &= \frac{P(a \wedge e)}{P(e)} \approx \frac{N_{tt}}{N_{tt} + N_{ft}} \\ &= \frac{0.1 \times 0.4 \times N}{0.1 \times 0.4 \times N + 0.3 \times 0.6 \times N} = 0.182 \end{aligned}$$

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$

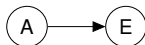


$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

- we can do better since we can **weight** the samples
- **weights = prob. that the evidence is observed**
- $N_t$  samples with  $A = \text{true}$  have weight of  $w_t = 0.1$   
this is  $P(E = \text{true} | A = \text{true})$
- $N_f$  samples with  $A = \text{false}$  have weight of  $w_f = 0.3$   
this is  $P(E = \text{true} | A = \text{false})$
- can do better because we don't need to generate the 90% of samples (when  $A = \text{true}$ ) that don't agree with the evidence - we simply assign all samples a weight of 0.1
- thus, we are **mixing exact inference** (the 0.1) with **sampling**.

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$



$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

- Compute sum of all weights of the samples with  $A = \text{true}$

$$W_t = \sum_i w_t \delta(a_i = \text{true}) = N_t \times 0.1$$

- Compute sum of all weights of the samples with  $A = \text{false}$

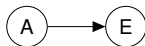
$$W_f = \sum_i w_f \delta(a_i = \text{false}) = N_f \times 0.3$$

- finally, compute

$$P(a|e) = \frac{W_t}{W_t + W_f} = \frac{0.1 \times 0.4 \times N}{0.1 \times 0.4 \times N + 0.3 \times 0.6 \times N}$$

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$

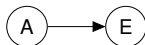


<i>A</i>	<i>P(E = true)</i>
<i>true</i>	0.1
<i>false</i>	0.3

- In fact, the *As* don't need to even be sampled from  $P(A)$
- Can be **sampled from some  $q(A)$** , say  $q(A = \text{true}) = 0.5$
- and each sample will have a new weight  $P(a)/q(a)$
- $q(A)$  is a **proposal** distribution.

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$

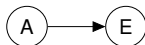


<i>A</i>	<i>P(E = true)</i>
<i>true</i>	0.1
<i>false</i>	0.3

- helps when it is hard to sample from  $P(A)$ , but we can evaluate  $P^*(A) \propto P(A)$  given a sample
- rejection sampling uses  $q = P$
- rejection sampling uses all variables including observed ones, and all weights on samples are set to 1.0

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$

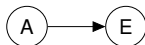


$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

- $N'_t = q(a)N$  samples with  $A = \text{true}$  have weight of  $0.1 \times \frac{P^*(a)}{q(a)} = 0.1 \times \frac{\alpha 0.4}{0.5}$
- $N'_f = q(\bar{a})N$  samples with  $A = \text{false}$  have weight of  $0.3 \times \frac{P^*(\bar{a})}{q(\bar{a})} = 0.3 \times \frac{\alpha 0.6}{0.5}$

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$



$A$	$P(E = \text{true})$
<i>true</i>	0.1
<i>false</i>	0.3

- total weight of all samples with  $A = \text{true}$

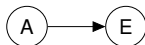
$$\begin{aligned} W'_t &= \sum_i w_i \delta(a_i = \text{true}) = N'_t \times 0.1 \times \frac{\alpha 0.4}{0.5} \\ &= 0.5N \times 0.1 \times \frac{\alpha 0.4}{0.5} = 0.1 \times \alpha \times 0.4 \times N \end{aligned}$$

- total weight of all samples with  $A = \text{false}$

$$\begin{aligned} W'_f &= \sum_i w_i \delta(a_i = \text{true}) = N'_f \times 0.3 \times \frac{\alpha 0.6}{0.5} \\ &= 0.5N \times 0.3 \times \frac{\alpha 0.6}{0.5} = 0.3 \times \alpha \times 0.6 \times N \end{aligned}$$

# Importance weights

$$\frac{P(A = \text{true})}{0.4}$$



<i>A</i>	<i>P(E = true)</i>
<i>true</i>	0.1
<i>false</i>	0.3

- finally, compute

$$P(a|e) = \frac{W'_t}{W'_t + W'_f} = \frac{0.1 \times \alpha \times 0.4 \times N}{0.1 \times \alpha \times 0.4 \times N + 0.3 \times \alpha \times 0.6 \times N}$$



# Example Proposal Distributions

- Sometimes we may want to choose a proposal distribution that is **different** than the actual probability distribution
- We may want to **skew** the proposal - because we may have some **additional knowledge** about the data, for example
- or, we can generate proposals **from the data itself** using some procedural knowledge that is not directly encoded in the BN
- Can be important in multiple/many dimensions,

# Stochastic sampling

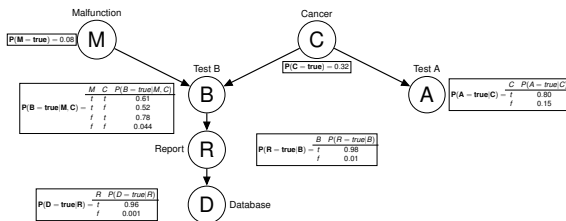
Recall variable elimination: To compute  $P(Z, Y_1 = v_1, \dots, Y_j = v_j)$ , we sum out the other variables,  $Z_1, \dots, Z_k = \{X_1, \dots, X_n\} - \{Z\} - \{Y_1, \dots, Y_j\}$ .

$$\begin{aligned} P(Z, Y_1 = v_1, \dots, Y_j = v_j) \\ = \sum_{Z_k} \cdots \sum_{Z_1} \prod_{i=1}^n P(X_i | \text{parents}(X_i))_{Y_1 = v_1, \dots, Y_j = v_j} \end{aligned}$$

Now, we sample  $Z_{l+1}, \dots, Z_k$  and sum  $Z_1, \dots, Z_l$ ,

$$= \sum_{s_j = \{z_{l+1,i}, \dots, z_{k,i}\}} \left[ \sum_{Z_1 \dots Z_l} \prod_{i=1}^l P(Z_i | \text{parents}(Z_i))_{Y_1 = v_1, \dots, Y_j = v_j} \right] \frac{P(Z_{l+1,i}, \dots, Z_{k,i})}{q(z_{l+1,i}, \dots, z_{k,i})}$$

# Importance Sampling example



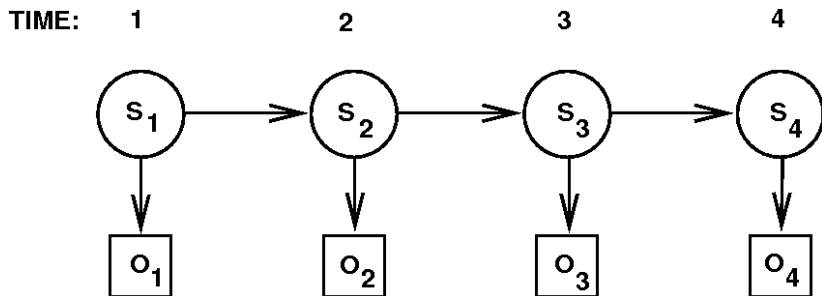
Compute  $P(B | D = \text{true}, A = \text{false})$  by sampling  $C$  and  $M$ .

- use  $q(C = \text{true}) = P(C = \text{true}) = 0.32$   
and  $q(M = \text{true}) = P(M = \text{true}) = 0.08$
- use  $q(C = \text{true}) = 0.5$   
and  $q(M = \text{true}) = P(M = \text{true}) = 0.08$
- use  $q(C = \text{true}) = q(M = \text{true}) = 0.5$

$$P(B | D = \text{true}, A = \text{false}) \propto \sum_{s_i = \{c_i, m_i\}} P(B, D = \text{true}, A = \text{false} | c_i, m_i)$$

see `sampling-inference.pdf`

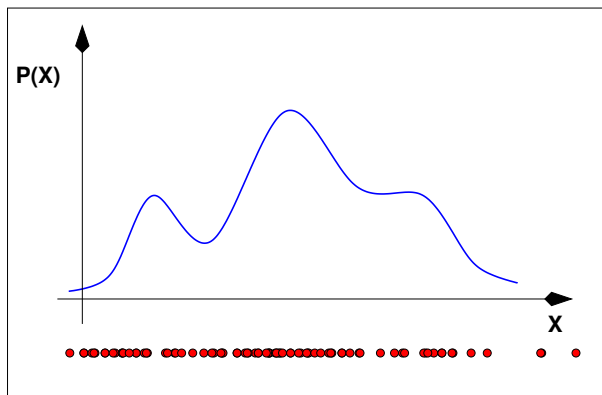
# Stochastic Sampling for HMMs (and other DBNS)



## Sequential Monte Carlo or Particle Filter

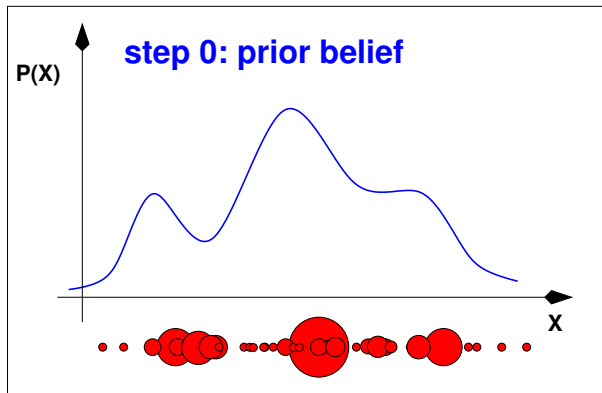
- sequential stochastic sampling
- keep track of  $P(S_t)$  at the current time  $t$
- represent  $P(S_t)$  with a set of samples
- update as new observations  $o_{t+1}$  arrive
  1. predict  $P(S_{t+1}) \propto P(S_{t+1}|S_t)$
  2. compute weights as  $P(o_{t+1}|S_{t+1})$
  3. resample according to weights

# Particle Filtering



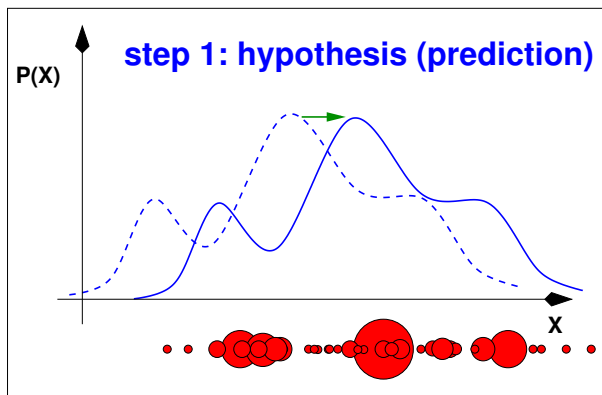
sample  $i$ :  $\{x_i\}$

# Particle Filtering



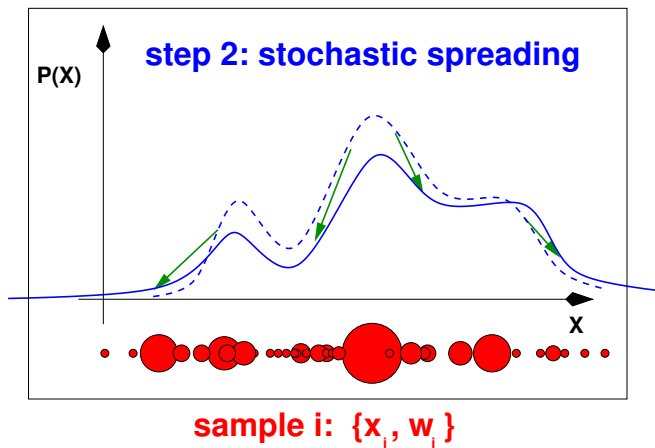
**sample i:  $\{x_i, w_i\}$**

# Particle Filtering



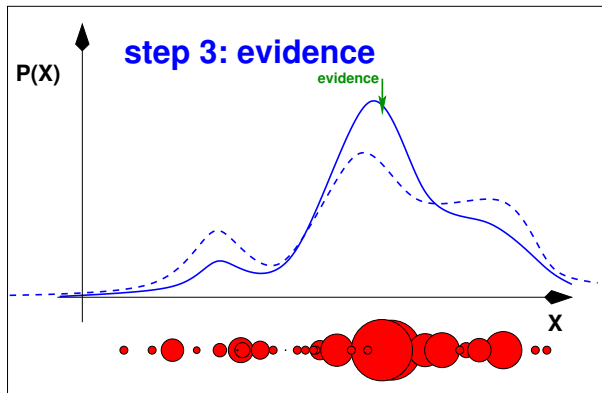
**sample  $i$ :  $\{x_i, w_i\}$**

# Particle Filtering



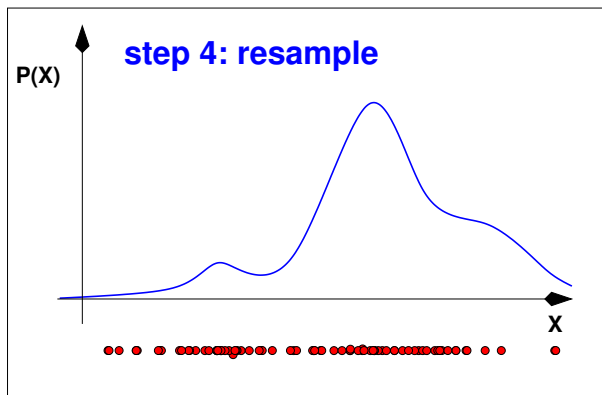


# Particle Filtering



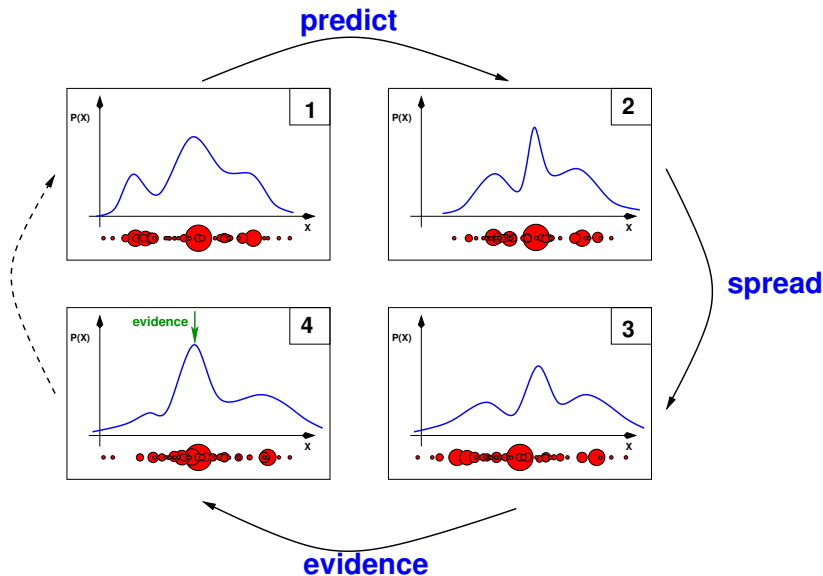
**sample  $i$ :  $\{x_i, w_i\}$**

# Particle Filtering

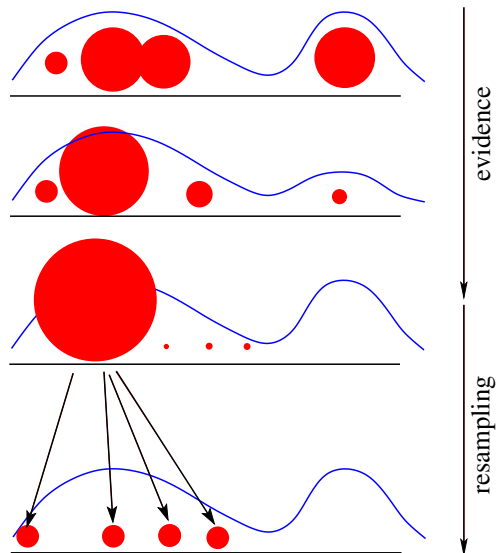


**sample  $i$ :  $\{x_i, w_i\}$**

# Bayesian Sequential Updates



# Resampling



- avoids degeneracies in the samples
- all importance weights  $\rightarrow 0$  except one
- performance of the algorithm depends on the resampling method

## Next:

- Supervised Learning under Uncertainty (Poole & Mackworth (2nd ed.)10.1,10.4)