

# Lecture 9c - Unsupervised Learning under Uncertainty

Jesse Hoey  
School of Computer Science  
University of Waterloo

June 29, 2022

Readings: Poole & Mackworth (2nd. Ed.) Chapt. 10.2,10.3,10.5

# Incomplete Data

- So far:
  - ▶ values of all attributes are known
  - ▶ learning is easy
- But many real-world problems have **hidden** variables (aka **latent** variables)
  - ▶ **Incomplete data**
  - ▶ Values of some attributes missing
- Incomplete data → **unsupervised learning**

# Maximum Likelihood learning

Recall: ML learning of Bayes net parameters for each variable  $V$  with parents  $pa(V)$ , and each value those parents can take on  $pa(V) = v$ :

$$\theta_{V=true, pa(V)=v} = P(V = true | pa(V) = v)$$

so that the ML learning of  $\theta$  is:

$$\theta_{V=true, pa(V)=v} = \frac{\text{number with } (V = true \wedge pa(V) = v)}{\text{number with } pa(V) = v}$$

Can add pseudocounts as priors

But what if some variable values are missing?

# Complete vs. Missing Data

For Cancer diagnosis example:

- Complete data (what we used to learn from in lecture 9a)

id	Malfnction	Cancer	TestB	TestA	Report	Database
1	false	false	true	true	false	false
2	false	true	true	true	true	true
3	false	true	true	true	true	true
4	false	false	false	true	false	false

...

- Incomplete (missing) data (more realistic)

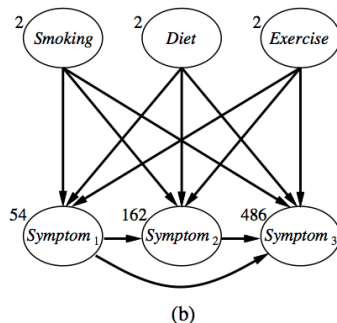
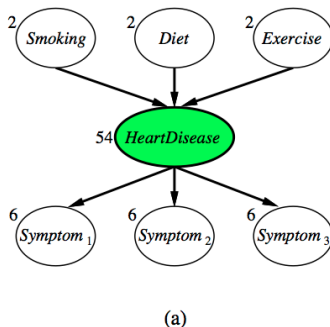
id	Malfnction	Cancer	TestB	TestA	Report	Database
1	?	?	?	true	?	false
2	?	?	?	true	?	false
3	true	?	?	true	?	false
4	?	true	?	true	?	false

...

# How to deal with missing data

## 1. Ignore hidden variables

number of parameters shown (variables have 3 values):



## 2. Ignore records with missing values

- ▶ does not work with true **latent** variables (e.g. always missing)


# Missing Data

- You **cannot ignore** missing data unless you know it is **missing at random.**
- Often data is missing because of something **correlated with a variable of interest.**
- For example: data in a clinical trial to test a drug may be missing because:
  - ▶ the patient dies,
  - ▶ the patient dropped out because of severe side effects,
  - ▶ they dropped out because they were better, or
  - ▶ the patient had to visit a sick relative.

— ignoring some of these may make the drug look better or worse than it is.
- In general you need to **model why data is missing.**

# Survivorship Bias

Bullet holes on planes returning from battle:  
where should the extra armour be installed?



Abraham Wald (WWII)

## “Direct” maximum likelihood

### 3. maximize likelihood directly

Suppose  $Z$  is hidden and  $E$  is observable, with values  $e$

$$\begin{aligned}h_{ML} &= \arg \max_h P(e|h) \\&= \arg \max_h \left[ \sum_Z P(e, Z|h) \right] \\&= \arg \max_h \left[ \sum_Z \prod_{i=1}^n P(X_i | \text{parents}(X_i), h)_{E=e} \right] \\&= \arg \max_h \left[ \log \sum_Z \prod_{i=1}^n P(X_i | \text{parents}(X_i), h)_{E=e} \right]\end{aligned}$$

Problem: can't push log inside the sum to linearize!



# Expectation-Maximization Algorithm

4. If we knew the missing values, computing  $h_{ML}$  would be easy again!

## Expectation-Maximization (EM):

A). **Guess**  $h_{ML}$

B). iterate:

- ▶ **expectation**: based on  $h_{ML}$ , compute expectation of missing values  $P(Z|h_{ML}, e)$
- ▶ **maximization**: based on expected missing values, compute new estimate of  $h_{ML}$

5. Really simple version (e.g. K-means algorithm):

- ▶ **expectation**: based on  $h_{ML}$ , compute **most likely** missing values  $\arg \max_Z P(Z|h_{ML}, e)$
- ▶ **maximization**: based on those missing values, you now have complete data, so compute new estimate of  $h_{ML}$  using ML learning as before

# $k$ -means algorithm

$k$ -means algorithm can be used for clustering :  
dataset of observables with input features  $X$  generated by one of a set of classes,  $C$  (e.g. Naïve Bayes,  $C \rightarrow X$ )

Inputs:

- training examples
- the number of classes,  $k$

Outputs:

- a representative value for each input feature for each class
- an assignment of examples to classes

Algorithm:

1. pick  $k$  means in  $X$ , one per class,  $C$
2. iterate until means stop changing:
  - a assign examples to  $k$  classes (e.g. as closest to current means)
  - b re-estimate  $k$  means based on assignment

# Expectation Maximization

- Approximate the maximum likelihood
- Start with a guess  $h_0$
- Iteratively compute:

$$h_{i+1} = \arg \max_h \sum_Z P(Z|h_i, e) \log P(e, Z|h)$$

- expectation: compute  $P(Z|h_i, e)$  ( "fills in" missing data )
- maximization: find new  $h$  that maximizes  $\sum_Z P(Z|h_i, e) \log P(e, Z|h)$
- can show that  $P(e|h_{i+1}) \geq P(e|h_i)$  when computed like this

# Expectation Maximization

Can show that:

$$\log P(e|h) \geq \sum_Z P(Z|e, h) \log P(e, Z|h)$$


EM finds a **local maximum** of right side: lower bound on left side  
log inside sum can **linearize the product**

$$\begin{aligned} h_{i+1} &= \arg \max_h \sum_Z P(Z|h_i, e) \log P(e, Z|h) \\ &= \arg \max_h \sum_Z P(Z|h_i, e) \log \prod_{j=1}^n P(X_j | \text{parents}(X_j), h)_{E=e} \\ &= \arg \max_h \sum_Z P(Z|h_i, e) \sum_{j=1}^n \log P(X_j | \text{parents}(X_j), h)_{E=e} \end{aligned}$$

EM **monotonically improves the likelihood**

$$P(e|h_{i+1}) \geq P(e|h_i)$$

# Naive Bayes with 4 input features



- Suppose  $k = 3$ , and  $\text{dom}(C) = \{1, 2, 3\}$ .
- $P(C|X_1, X_2, X_3, X_4) \propto P(X_1 \dots X_4|C)P(C)$
- can be computed if we know  $P(C)$  and  $P(X_i|C)$
- **EM idea**: based on current  $P(C)$  and  $P(X_i|C)$ , compute  $P(C|X_1 \dots X_4) \forall C \in \{1, 2, 3\}$
- use  $P(C|X_1 \dots X_4)$  as **partial data** in ML learning

# Augmented Data Method — E step – Naive Bayes

$$P(C = 1|X_1 = t, X_2 = f, X_3 = t, X_4 = t) = 0.4$$

$$P(C = 2|X_1 = t, X_2 = f, X_3 = t, X_4 = t) = 0.1$$

$$P(C = 3|X_1 = t, X_2 = f, X_3 = t, X_4 = t) = 0.5:$$

missing data (C)  $\longrightarrow$  filled in data

$X_1$	$X_2$	$X_3$	$X_4$		
$\vdots$	$\vdots$	$\vdots$	$\vdots$		
$t$	$f$	$t$	$t$		
$\vdots$	$\vdots$	$\vdots$	$\vdots$		

 $\longrightarrow$ 

$X_1$	$X_2$	$X_3$	$X_4$	$C$	$val$
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
$t$	$f$	$t$	$t$	1	0.4
$t$	$f$	$t$	$t$	2	0.1
$t$	$f$	$t$	$t$	3	0.5
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$

call this  $A[X_1, \dots, X_4, C]$

Compute the statistics for each feature and class:

$$M_i[X_i, C] = \sum_{X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_n} A[X_1, \dots, X_n, C]$$

$$M[C] = \sum_{X_i} M_i[X_i, C]$$

$M[C]$  is unnormalized marginal.

Compute the statistics for each feature and class:

$$M_i[X_i, C] = \sum_{X_1, \dots, X_{i-1}, X_{i+1}, \dots, X_n} A[X_1, \dots, X_n, C]$$

$$M[C] = \sum_{X_i} M_i[X_i, C]$$

$M[C]$  is unnormalized marginal.

Compute probabilities by normalizing:

$$P(X_i|C) = M_i[X_i, C]/M[C]$$

$$P(C) = M[C]/s$$

Pseudo-counts can also be added.



# General Bayes Network EM

**Complete data**: Bayes Net Maximum Likelihood

$$\theta_{V=true, pa(V)=v} = \frac{\text{number in } e \text{ with } (V = true \wedge pa(V) = v)}{\text{number in } e \text{ with } pa(V) = v}$$

**Incomplete data**: Bayes Net Expectation Maximization

observed variables  $X$  and missing variables  $Z$

Start with some **guess** for  $\theta$ ,

**E Step**: Compute weights for each data  $x_i$  and latent variable(s) value(s)  $z_j$  (using e.g. variable elimination)

$$w_{ij} = P(z_j | \theta, x_i)$$

**M Step**: Update parameters:

$$\theta_{V=true, pa(V)=v} = \frac{\sum_{ij} w_{ij} | V = true \wedge pa(V) = v \text{ in } \{x_i, z_j\}}{\sum_{ij} w_{ij} | pa(V) = v \text{ in } \{x_i, z_j\}}$$

# Belief network structure learning (I)

$$P(model|data) = \frac{P(data|model) \times P(model)}{P(data)}.$$

- A model here is a belief network.
- A bigger network can always fit the data better.
- $P(model)$  lets us encode a preference for smaller networks (e.g., using the description length).
- You can search over network structure looking for the most likely model.

# Belief network structure learning (II)

- can do **independence tests** to determine which features should be the parents
- **XOR problem**: just because features do not give information individually, does not mean they will not give information in combination
- ideal: **Search over total orderings** of variables

## Next:

- Planning with uncertainty (Poole & Mackworth (2nd. Ed.) chapter 9.1-9.3,9.5)
- Reinforcement Learning (Poole & Mackworth (2nd. Ed.) chapter 12.1,12.3-12.9)