

Lecture 10 - Planning under Uncertainty (III)

Jesse Hoey
School of Computer Science
University of Waterloo

July 6, 2022

Readings: Poole & Mackworth (2nd ed.) Chapter 12.1, 12.3-12.9

What should an agent do given:

- **Prior knowledge** possible states of the world
possible actions
- **Observations** current state of world
immediate reward / punishment
- **Goal** act to maximize accumulated reward

Like decision-theoretic planning, except model of dynamics and model of reward not given.

- We assume there is a **sequence of experiences** :

state, action, reward, state, action, reward,

- What should the agent do next?
- It must decide whether to:
 - ▶ **explore** to gain more knowledge
 - ▶ **exploit** the knowledge it has already discovered

Reinforcement Learning: “Bandit” problem



Each machine has a $Pr(win)$... but you don't know what it is...

Which machine should you play?

Why is reinforcement learning hard?

- What actions are responsible for the reward may have occurred a **long time before** the reward was received.
- The **long-term effect** of an action of the robot depends on what it will do in the future.
- The **explore-exploit dilemma**: at each time should the robot be greedy or inquisitive?

Reinforcement learning: main approaches

- search through a space of policies (controllers)
- **Model Based RL**: **learn a model** consisting of state transition function $P(s'|a, s)$ and reward function $R(s, a, s')$; **solve** this as an MDP.
- **Model-Free RL** **learn $Q^*(s, a)$** , use this to guide action.

Temporal Differences

- Suppose we have a sequence of values:

$$v_1, v_2, v_3, \dots$$

And want a **running estimate** of the average of the first k values:

$$A_k = \frac{v_1 + \dots + v_k}{k}$$

Temporal Differences (cont)

- When a new value v_k arrives:

$$A_k = \frac{v_1 + \cdots + v_{k-1} + v_k}{k}$$

$$\begin{aligned} kA_k &= v_1 + \cdots + v_{k-1} + v_k \\ &= (k-1)A_{k-1} + v_k \end{aligned}$$

$$A_k = \frac{k-1}{k}A_{k-1} + \frac{1}{k}v_k$$

Let $\alpha = \frac{1}{k}$, then

$$\begin{aligned} A_k &= (1 - \alpha)A_{k-1} + \alpha v_k \\ &= A_{k-1} + \alpha(v_k - A_{k-1}) \end{aligned}$$

“TD formula”

- Often we use this update with α fixed.

- **Idea:** store $Q[\text{State}, \text{Action}]$; update this as in **asynchronous value iteration**, but using experience (empirical probabilities and rewards).
- Suppose the agent has an experience $\langle s, a, r, s' \rangle$
- This provides one piece of data to update $Q[s, a]$.
- The experience $\langle s, a, r, s' \rangle$ provides the data point:

$$r + \gamma \max_{a'} Q[s', a']$$

which can be used in the **TD formula** giving:

$$Q[s, a] \leftarrow Q[s, a] + \alpha \left(r + \gamma \max_{a'} Q[s', a'] - Q[s, a] \right)$$

begin

initialize $Q[S, A]$ arbitrarily

observe current state s

repeat forever:

select and carry out an action a

observe reward r and state s'

$Q[s, a] \leftarrow Q[s, a] + \alpha (r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

$s \leftarrow s'$;

end-repeat

end

Properties of Q-learning

- Q-learning **converges to the optimal policy**, no matter what the agent does, as long as it **tries each action in each state enough (infinitely often)**.
- But what should the agent do?
 - ▶ **exploit**: when in state s , select the action that maximizes $Q[s, a]$
 - ▶ **explore**: select another action

Exploration Strategies

- The **ϵ -greedy** strategy: choose a random action with probability ϵ and choose a best action with probability $1 - \epsilon$.
- **Softmax** action selection: in state s , choose action a with probability

$$\frac{e^{Q[s,a]/\tau}}{\sum_a e^{Q[s,a]/\tau}}$$

where $\tau > 0$ is the **temperature**. Good actions are chosen more often than bad actions; τ defines how often good actions are chosen. For $\tau \rightarrow \infty$, all actions are equiprobable. For $\tau \rightarrow 0$, only the best is chosen.


Exploration Strategies

- **optimism in the face of uncertainty**: initialize Q to values that encourage exploration.
- **Upper Confidence Bound (UCB)**: Also store $N[s, a]$ (number of times that state-action pair has been tried) and use

$$\arg \max_a \left[Q(s, a) + k \sqrt{\frac{N[s]}{N[s, a]}} \right]$$

where $N[s] = \sum_a N[s, a]$


Example: studentbot



state variables:

- **tired**: studentbot is tired (no/a bit/very)
- **passtest**: studentbot passes test (no/yes)
- **knows**: studentbot's state of knowledge (nothing/a bit/a lot/everything)
- **goodtime**: studentbot has a good time (no/yes)


Example: studentbot



studentbot actions:

- **study**: studentbot's knowledge increases, studentbot gets tired
- **sleep**: studentbot gets less tired
- **party**: studentbot has a good time, but gets tired and loses knowledge
- **take test**: studentbot takes a test (can take test anytime)

Example: studentbot

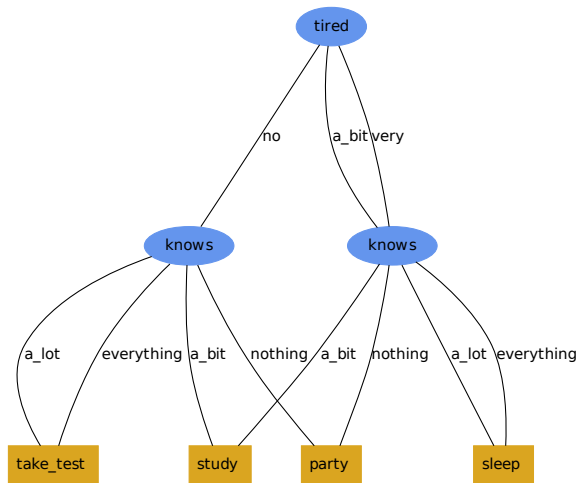


studentbot rewards:

- **+20** if studentbot passes the test
- **+2** if studentbot has a good time when not very tired

basic tradeoff: short term vs. long-term rewards

Studentbot Policy



Model-based Reinforcement Learning

- Model-based reinforcement learning uses the experiences in a more effective manner.
- It is used when collecting experiences is expensive (e.g., in a robot or an online game), and you can do lots of computation between each experience.
- **Idea** : learn the MDP and interleave acting and planning.
- After each experience, update probabilities and the reward, then do some steps of **asynchronous value iteration**.

Model-based Reinforcement Learning

- Model-based reinforcement learning uses the experiences in a more effective manner.
- It is used when collecting experiences is expensive (e.g., in a robot or an online game), and you can do lots of computation between each experience.
- Idea: learn the MDP and interleave acting and planning.
- After each experience, update probabilities and the reward, then do some steps of asynchronous value iteration.

Model-based learner

Data Structures: $Q[S, A]$, $T[S, A, S]$, $R[S, A]$

Assign Q , R arbitrarily, T = prior counts

α is learning rate

observe current state s

repeat forever:

 select and carry out action a

 observe reward r and state s'

$T[s, a, s'] \leftarrow T[s, a, s'] + 1$

$R[s, a] \leftarrow \alpha \times r + (1 - \alpha) \times R[s, a]$

repeat for a while (asynchronous VI):

 select state s_1 , action a_1

 let $P = \sum_{s_2} T[s_1, a_1, s_2]$

$Q[s_1, a_1] \leftarrow \sum_{s_2} \frac{T[s_1, a_1, s_2]}{P} \left(R[s_1, a_1] + \gamma \max_{a_2} Q[s_2, a_2] \right)$

$s \leftarrow s'$

Off/On-policy Learning

- Q-learning does **off-policy** learning: it learns the value of the optimal policy, no matter what it does.
- This could be bad if the exploration policy is dangerous.
- **On-policy** learning learns the value of the policy being followed.
e.g., act greedily 80% of the time and act randomly 20% of the time
- If the agent is actually going to explore, it may be better to optimize the actual policy it is going to do.
- **SARSA** uses the experience $\langle s, a, r, s', a' \rangle$ to update $Q[s, a]$.

begin

initialize $Q[S, A]$ arbitrarily

observe current state s

select action a using a policy based on Q

repeat forever:

 carry out an action a

 observe reward r and state s'

 select action a' using a policy based on Q

$Q[s, a] \leftarrow Q[s, a] + \alpha (r + \gamma Q[s', a'] - Q[s, a])$

$s \leftarrow s'$;


$a \leftarrow a'$;

end-repeat


end

Large State Spaces

- **Computer Go**: 3^{361} states



- **Atari** Games $210 \times 160 \times 3$ dimensions (pixels)



Q-function Approximations

- Let $s = (x_1, x_2, \dots, x_N)^T$

- **Linear**

$$Q_w(s, a) \approx \sum_i w_{ai} x_i$$

- **Non-linear** (e.g. neural network)

$$Q_w(s, a) \approx g(x; w)$$

Recall: Logistic Regression

Logistic function of linear weighted inputs:

$$\hat{Y}^{\bar{w}}(e) = f(w_0 + w_1 X_1(e) + \dots + w_n X_n(e)) = f\left(\sum_{i=0}^n w_i X_i(e)\right)$$

The sum of squares error is:

$$Error(E, \bar{w}) = \sum_{e \in E} \left[Y(e) - f\left(\sum_{i=0}^n w_i * X_i(e)\right) \right]^2$$

The partial derivative with respect to weight w_i is:

$$\frac{\partial Error(E, \bar{w})}{\partial w_i} = -2 * \delta * f' \left(\sum_i w_i * X_i(e) \right) * X_i(e)$$

where $\delta = (Y(e) - f(\sum_{i=0}^n w_i X_i(e)))$.

Thus, each example e updates each weight w_i by

$$w_i \leftarrow w_i + \eta * \delta * f' \left(\sum_i w_i * X_i(e) \right) * X_i(e)$$

Approximating the Q-function

- for experience tuple s, a, r, s' we have:
 - ▶ target Q-function: $R(s) + \gamma \max_{a'} Q_w(s', a')$ or $R(s) + \gamma Q_w(s', a')$
 - ▶ current Q-function: $Q_w(s, a)$
- Squared error:

$$Err(w) = \frac{1}{2} \left[Q_w(s, a) - R(s) - \gamma \max_{a'} Q_w(s', a') \right]^2$$

- Gradient:

$$\frac{\partial Err}{\partial w} = \left[Q_w(s, a) - R(s) - \gamma \max_{a'} Q_w(s', a') \right] \frac{\partial Q_{w(s,a)}}{\partial w}$$

SARSA with linear function approximation

Given γ :discount factor; α :learning rate

Assign weights $\bar{w} = \langle w_0, \dots, w_n \rangle$ arbitrarily

begin

 observe current state s

 select action a

repeat forever:

 carry out action a

 observe reward r and state s'

 select action a' (using a policy based on Q_w)

 let $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

 For $i = 0$ to n

$$w_i \leftarrow w_i + \alpha \times \delta \times \frac{\partial Q_w(s, a)}{\partial w}$$

$s \leftarrow s'; a \leftarrow a';$

end-repeat

end

- Linear Q-learning ($Q_w(s, a) \approx \sum_i w_{ai}x_i$) converges under same conditions as Q-learning

$$w_i \leftarrow w_i + \alpha [Q_w(s, a) - R(s) - \gamma Q_w(s', a')] x_i$$

- Non-linear Q-learning (e.g. neural network, $Q_w(s, a) \approx g(x; w)$) may diverge
 - ▶ Adjusting w to increase Q at (s, a) might introduce errors at nearby state-action pairs.

Mitigating Divergence

Two tricks used in practice:

1. Experience Replay
2. Use two Q function (two networks):
 - ▶ Q network (currently being updated)
 - ▶ Target network (occasionally updated)

Experience Replay

- **Idea:** Store previous experiences (s, a, r, s', a') in a buffer and sample a mini-batch of previous experiences at each step to learn by Q-learning
- **Breaks correlations** between successive updates (more stable learning)
- Few interactions with environment needed to converge (greater **data efficiency**)

Target Network

- **Idea**: use a separate target network that is **updated only periodically**
- target network has weights \bar{w} and computes $Q_{\bar{w}}(s, a)$
- repeat for each (s, a, r, s', a') in **mini-batch**:

$$w \leftarrow w + \alpha \left[Q_w(s, a) - R(s) - \gamma Q_{\bar{w}}(s', a') \right] \frac{\partial Q_w(s, a)}{\partial w}$$

- $\bar{w} \leftarrow w$

Deep Q-Network

Assign weights $\bar{w} = \langle w_0, \dots, w_n \rangle$ at random in $[-1, 1]$

begin

observe current state s

select action a

repeat forever:

carry out action a

observe reward r and state s'

select action a' (using a policy based on Q_w)

add (s, a, r, s', a') to experience buffer

Sample mini-batch of experiences from buffer

For each experience $(\hat{s}, \hat{a}, \hat{r}, \hat{s}', \hat{a}')$ in mini-batch:

let $\delta = \hat{r} + \gamma Q_{\bar{w}}(\hat{s}', \hat{a}') - Q_w(\hat{s}, \hat{a})$

$w \leftarrow w + \alpha \times \delta \times \frac{\partial Q_w(\hat{s}, \hat{a})}{\partial w}$

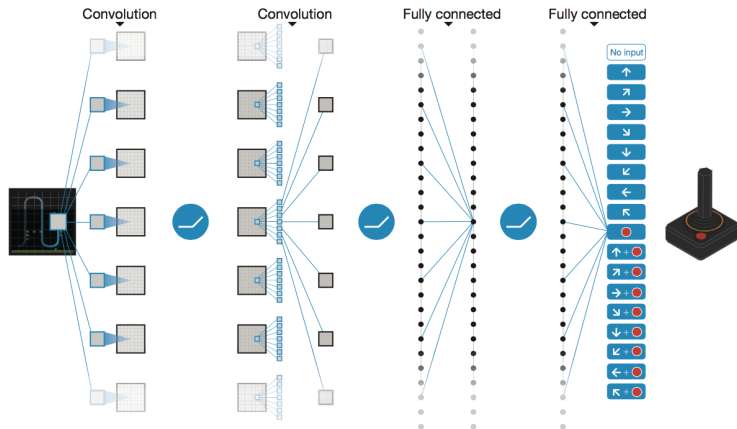
$s \leftarrow s'; a \leftarrow a';$

every c steps, update target $\bar{w} \leftarrow w$

end-repeat


end

Deep Q-Network for Atari



from: Mnih *et. al.* Human-level control through deep reinforcement learning.
Nature 18(7540):529–533 2015.

Deep Q-Network vs. Linear Approx.



Bayesian Reinforcement Learning

- Include the parameters (transition function and observation function) in the state space
- Model-based learning though inference (belief state)
- State space is now continuous,
belief space is a space of continuous functions
- Can mitigate complexity by modeling reachable beliefs
- optimal exploration-exploitation tradeoff.

Next:

- Recap